

# Overview on Optimization Algorithms

Markus A. Köbis

NTNU, TMA4180, Optimering I, Spring 2021

2021-02-26

# A Review Session

## Idea

- OPTION A - a live session with a presentation highlighting the 'big picture' of what we have learned so far

## What we don't do today

- Learn new material ;-)
- Go through proofs

Questions are welcome!

# Talking about the ‘Big Picture’(I)

## Topics of Today:

Vurderingsform	Vektning	Varighet	Hjelpemidler
Arbeider	30/100		
Hjemmeeksamen	70/100	4 timer	

### Faglig innhold

Første og andre ordens nødvendige og tilstrekkelige (Karush-Kuhn-Tucker) optimalitetsbetingelser for ubegrenede og begrenede optimeringsproblemer i endelig-dimensjonale vektorrom. Grunnleggende konveksanalyse og Lagranges dualitetsteori og deres anvendelser for optimeringsproblemer og algoritmer. Oversikt over moderne optimeringsteknikker og algoritmer for glatte problemer (inklusive line-search/trust-region, kvari-Newton, indre punkt og aktive sett metoder, SQP og augmented Lagrangian teknikker). Grunnleggende metoder for derivat-fri og ikke-glatte optimeringsproblem.

### Læringsutbytte

Studenten som møter læringsmålene for kurset skal kunne:

- (i) vurdere eksistens og entydighet av løsninger til et gitt optimeringsproblem;
- (ii) validere konveksitet av funksjoner, sett, og optimeringsproblemer;
- (iii) utlede nødvendige og tilstrekkelige optimalitetsbetingelser for et gitt optimeringsproblem;
- (iv) løse små optimeringsproblemer analytisk;
- (v) forklare de underliggende prinsipper og begrensninger av moderne teknikker og algoritmer for optimering;
- (vi) anslå konvergenshastigheten og kompleksitetskrav i ulike optimeringsalgoritmer;
- (vii) implementere optimeringsalgoritmer på en datamaskin;
- (viii) bruke optimeringsalgoritmer for å løse modellproblemer i ingenør- og realfag.

### Læringsformer og aktiviteter

## Talking about the 'Big Picture' (II)

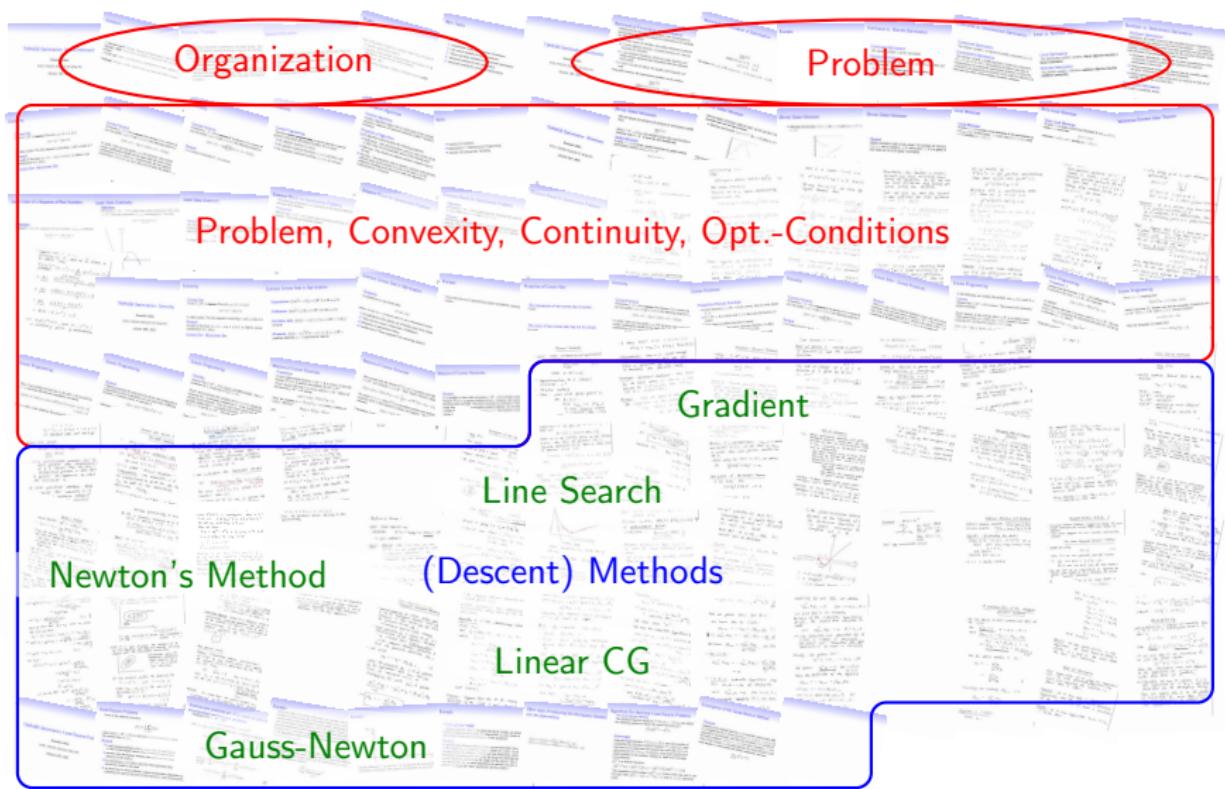
Topics of Today:

Recap of Theory and Methods for **Unconstrained** Optimization Problems

$$\min_{x \in \mathbb{R}^n} f(x)$$

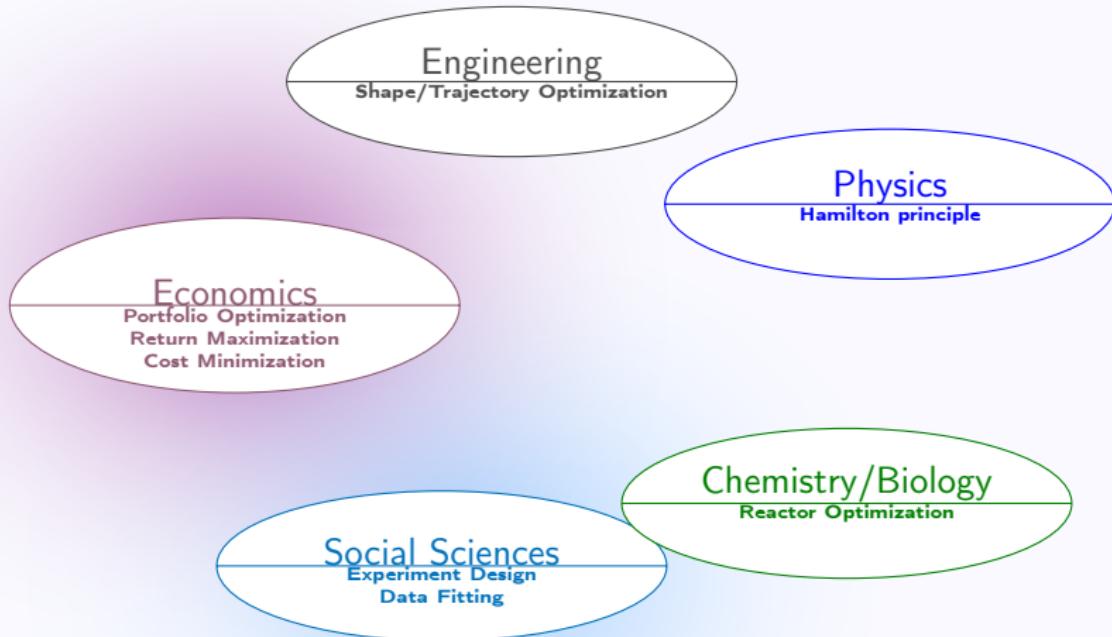
- (Practical) Problems and Challenges
- Theoretical Considerations
- **Algorithms**

## Talking about the 'Big Picture' (III)



# Problems and Challenges (I)

## Optimization Problems in Science and Engineering



# Taxonomy of Optimization Problems (Not exhaustive)



Source: <https://neos-guide.org/content/optimization-taxonomy>

# Quality of Algorithms

Properties of 'good' algorithms (Why isn't there a best one?)

- Accuracy
- Efficiency (also: Scalability)
  - Speed
    - Total execution time (avoid NP if possible, parallel vs. sequential)
    - Alternative: approximation quality after fixed time
    - Number of computer operations
      - Core function/ Linear algebra/GPU-routine-calls
      - Floating point operations (plus, times, roots, exp./trig, ...)
      - Bit operations
      - Number of  $f/\nabla f/\nabla^2 f$  calls
    - Memory
  - Robustness
    - Reliably find all solutions for a variety of problems
    - Feedback if algorithm unsuccessful, reproducibility
    - Provide additional information on problem structure
  - User-friendliness, ability to include expert knowledge, interactivity
  - Easily extendable and maintainable
  - Cheap, trusted, supported, ...

## Building the 'Big Picture'

Mathematical Foundations  $\leftrightarrow$  Algorithms

# When Do Minimizers Exist?

## Weierstrass' theorem

### Weierstrass Theorem

A **continuous** function on a **compact** domain attains its minimum.

- In practice: Very often enough
- Can one weaken the 'continuity' requirement?  
e.g. complicated functions in technical applications
- Can one weaken the 'compactness' requirement?  
e.g. optimal control/problems in mathematical physics



Weierstrass

Karl Weierstrass, 1815–1897

Source: Public Domain,

<https://commons.wikimedia.org/w/index.php?curid=324146>

# Convexity

## Convex Set

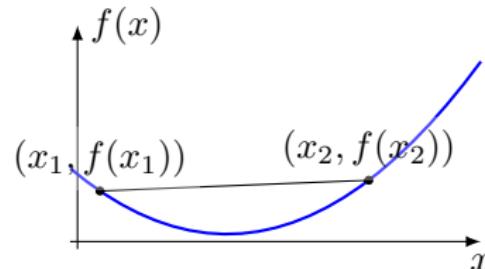
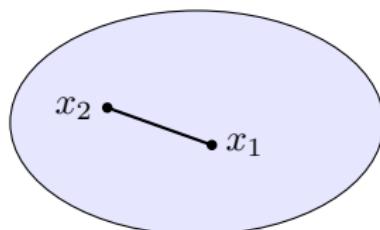
A set  $S \subseteq \mathbb{R}^n$  is **convex** if for all  $x_1, x_2 \in S$ ,  $\lambda \in [0, 1]$

$$\underbrace{\lambda x_1 + (1 - \lambda)x_2}_{\text{convex combination}} \in S.$$

## Convex Function

A function  $f: S \rightarrow \mathbb{R}$  is **convex** if its domain  $S$  is a convex set and if for any two points  $x_1, x_2 \in S$  and  $\lambda \in [0, 1]$  the following property is satisfied:

$$f(\lambda x_1 + (1 - \lambda)x_2) \leq \lambda f(x_1) + (1 - \lambda)f(x_2).$$

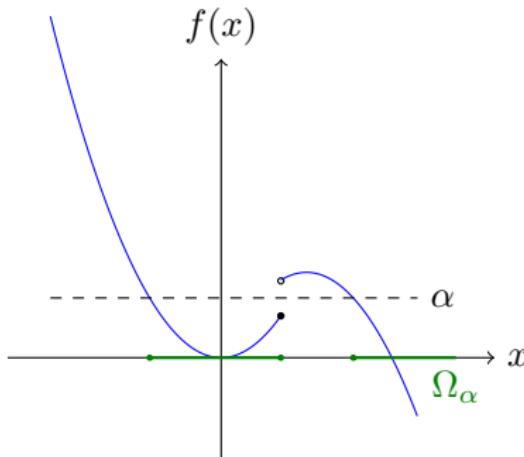


## (Lower Semi-) Continuity

### Definition

A function  $f: \mathbb{R}^n \rightarrow \mathbb{R}$  is called *lower semi-continuous*, if for every  $x \in \mathbb{R}^n$  and every sequence  $(x_k)_{k \in \mathbb{N}}$  converging to  $x$  we have

$$f(x) \leq \liminf_{k \rightarrow \infty} f(x_k).$$



### Remark

An alternative (equivalent) definition of lower semi-continuity is the following: A function  $f: \mathbb{R}^n \rightarrow \mathbb{R}$  is lower semi-continuous, if the *lower level set*

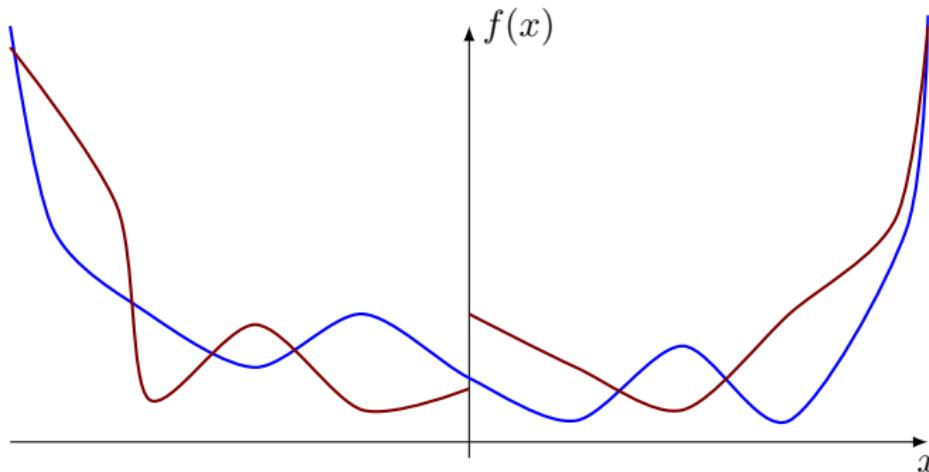
$$\Omega_\alpha(f) := \{x \in \mathbb{R}^n : f(x) \leq \alpha\}$$

is closed for every  $\alpha \in \mathbb{R}$ .

# Coercivity

## Definition

A function  $f: \mathbb{R}^n \rightarrow \mathbb{R}$  is called coercive, if we have for every sequence  $(x_k)_{k \in \mathbb{N}}$  with  $\|x_k\| \rightarrow \infty$  that  $f(x_k) \rightarrow \infty$ .



## Weierstrass revisited

### Proposition

Assume that  $f: \mathbb{R}^n \rightarrow \mathbb{R}$  is lower semi-continuous and that  $\Omega \subset \mathbb{R}^n$  is compact. Then the optimization problem

$$\min_{x \in \Omega} f(x)$$

has at least one global minimizer.

### Theorem

Assume that the function  $f: \mathbb{R}^n \rightarrow \mathbb{R}$  is lower semi-continuous and coercive. Then the optimization problem  $\min_{x \in \mathbb{R}^n} f(x)$  admits at least one global minimizer.

# The Role of Convexity

## Theorem

If  $f: \mathbb{R}^n \rightarrow \mathbb{R}$  is convex, any local minimizer is a global minimizer. (This implies that optimality conditions (see below) can be used for global optimization.)

# Designing Algorithms

## Optimization Problem:

$$\min_{x \in \mathbb{R}^n} f(x)$$

### How to proceed?

Design computer method to find good points in  $\mathbb{R}^n$  until you are finished

- (a) Can the search be based on existing mathematical tools? (e.g. Solve (non-)linear systems, feasibility problem solvers, simulation methods etc.)
- (b) How can you know that the algorithm is 'finished'?

→ Optimality Conditions

# Optimality Conditions (I)

## First Order Necessary Condition

Let  $x^* \in \mathbb{R}^n$  be a *local* minimizer, and let  $f: \mathbb{R}^n \rightarrow \mathbb{R}$  be continuously differentiable in an open neighborhood of  $x^*$ . Then

$$\nabla f(x^*) = \mathbf{0}.$$

This condition classifies  $x^*$  as a stationary point

## Optimality Conditions (II)

### Second Order Necessary Condition

Let  $x^* \in \mathbb{R}^n$  be a local minimizer of  $f: \mathbb{R}^n \rightarrow \mathbb{R}$ . If  $\nabla^2 f$  exists and is continuous in an open neighborhood of  $x^*$ , then  $\nabla f(x^*) = \mathbf{0}$  and  $\nabla^2 f(x^*)$  is positive semi-definite.

### Second Order Sufficient Condition

Let  $\nabla^2 f$  be continuous in an open neighborhood of  $x^* \in \mathbb{R}^n$ ,  $\nabla f(x^*) = \mathbf{0}$  and let  $\nabla^2 f(x^*)$  be positive definite. Then  $x^*$  is a strict local minimizer.

### Geometry/Practice

- Properties of a **Quadratic Approximation on  $f$**
- In (large-scale) algorithms often not/only approximately used.

## Designing Algorithms (II)

$$\min_{x \in \mathbb{R}^n} f(x)$$

**Algorithm:** Computerized function (Input → output scheme) with limited internal memory, calculation power and precision.

**Picture:** Point cloud

## Designing Algorithms (III)

### Versatile methods for 'any' optimization problem

1. First Attempt: **Grid Search**: Evaluate  $f$  along a predefined grid and choose some good grids to refine this further.  
‘Curse of Dimension’: Number of grid points grows exponentially with dimension  $n$  of  $x$ .
2. Improvement (a): Solve a series of 1d-problems using ‘classical methods’  
How to define these?
3. Improvement (b): Algorithms based on Heuristics/Physics/Biology  
e.g. Nelder-Mead Simplex Method, Simulated Annealing, Evolutionary Algorithms
4. (Often) Better: **Iterative Procedure**: Choose a starting point  $x_0 \in \mathbb{R}^n$  and search along a sequence

$$x_0, x_1, x_2, \dots, x_k, x_{k+1}, \dots \rightarrow x^*$$

with improved properties of  $x_{k+1}$  over  $x_k$ .

Given the scarce information ( $f$  is just input→output), the only way to measure ‘improved properties’ is through function values → **descent methods**

scaling invariance, also important to account for easy extension and maintenance of code

# Descent Methods (I)

## Basic Idea of Descent Methods

1. Start with initial  $x_0 \in \mathbb{R}^n$

2. **For**  $k = 0, 1, 2, \dots$

Find  $x_{k+1}$  such that  $f(x_{k+1}) < f(x_k)$

If some stopping criterion is fulfilled: STOP

**Output:**  $x^* = x_{k+1}$

How to find a better  $x$ ? Define the search direction

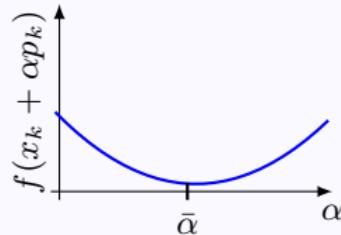
$$\lambda \cdot p_k := x_{k+1} - x_k \quad \Leftrightarrow \quad x_{k+1} = x_k + p_k$$

When is  $p_k$  good?

## Definition

For a given  $x_k \in \mathbb{R}^n$ , a direction (nonzero vector (of undefined length))  $p_k \in \mathbb{R}^n$  is called a **descent direction** if there exists  $\bar{\alpha} > 0$  such that

$$f(x_k + \alpha p_k) < f(x_k) \quad \forall \alpha \in (0, \bar{\alpha}).$$



## Descent Methods (II)

**Lemma: Characerizing descent directions for differentiable  $f$**

Let  $f: \mathbb{R}^n \rightarrow \mathbb{R}$  be continuously differentiable.

1. Any direction  $p$  satisfying

$$\nabla f(x)^\top \cdot p < 0$$

is a descent direction.

2. The negative gradient  $-\nabla f(x)$  is a descent direction (if not zero)
3. For any positive definite matrix  $B \in \mathbb{R}^{n \times n}$ , the direction

$$-B \cdot \nabla f(x)$$

is a descent direction (if  $x$  is not stationary).

## Line Search (I)

In general form of the descent algorithm:

$$x_{k+1} := x_k + \alpha_k \cdot p_k$$

What is a good  $\alpha_k$ ?

1. Constant  $\alpha$ :

- Good for some special algorithms and when problem structure is well-understood
- $\alpha$  too large: ‘overshooting’ and method may diverge
- $\alpha$  too small: Potentially many steps and slow convergence

2. Determining a perfect value, **exact line search**:  $\alpha_k := \operatorname{argmin}_x \underbrace{\Phi_k(\alpha)}_{=f(x_k + \alpha \cdot p_k)}$

- Good if  $\Phi_k$ -minimization is easy.

3. **Inexact Line Searches**: Use heuristics to obtain adequate step lengths at justifiable cost

# Line Search (I)

## Backtracking

### (Sub-) Algorithm: Backtracking

**Problem:** Find a good/quick approximation to the solution of  $\min_{\alpha} \Phi_k(\alpha)$ .

**Parameter:**  $\hat{\alpha} \in \mathbb{R}$  as initial guess, reduction constant  $\rho \in (0, 1)$

**Set:**  $\alpha = \hat{\alpha}$

**While:** some stopping criterion for  $\phi_k(\alpha)$  is not fulfilled:

**set:**  $\alpha = \rho \cdot \alpha$ .

### In Practice

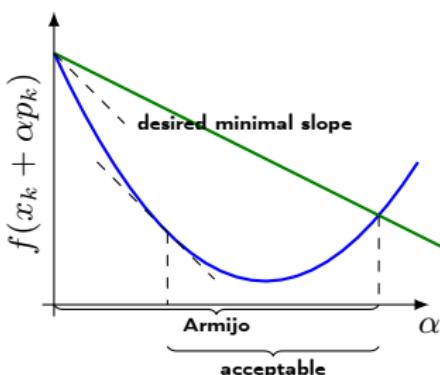
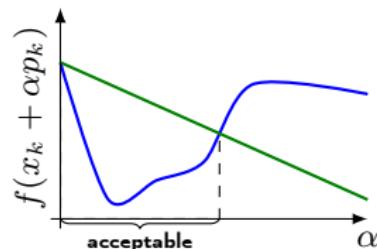
- Also set a maximum number of iterations
- Sometimes: Re-use  $\alpha$  from previous run of algorithm
- Improvements based on 1d equation solvers (regula falsi, interpolation etc.)
- $\rho \in (0.4, 0.7)$

## Line Search (II)

### Armijo (Sufficient Decrease) Condition

Given  $c_1 \in (0, 1)$ , we require

$$\underbrace{f(x_k + \alpha_k p_k)}_{\Phi_k(\alpha_k)} \leq \underbrace{f(x_k) + c_1 \cdot \alpha_k \cdot \nabla f(x_k)^\top \cdot p_k}_{\Phi_k(0) + \alpha_k \cdot c_1 \cdot \Phi'_k(0) =: l(\alpha_k)}$$



### Wolfe (Curvature) Condition(s)

Given, furthermore, a constant  $c_2 \in (c_1, 1)$ , we require additionally

$$\underbrace{\nabla f(x_k + \alpha_k \cdot p_k)^\top \cdot p_k}_{\Phi'_k(\alpha_k)} \geq \underbrace{c_2 \cdot \nabla f(x_k)^\top \cdot p_k}_{c_2 \cdot \Phi'_k(0)}$$

*Strong Wolfe condition:* Also restrict positive slope

Existence of such  $\alpha_k$  given if  $f$  smooth and bounded from below along the search line.

## Line Search (III)

Fundamental Result for Line Search Methods

### Theorem (Zoutendijk)

Consider any iteration of the form

$$x_{k+1} = x_k + \alpha_k \cdot p_k$$

where  $p_k$  is a descent direction and  $\alpha_k$  satisfies the Wolfe conditions for  $0 < c_1 < c_2 < 1$ . Suppose that  $f$  is bounded below, differentiable with L-continuous gradient on any open set containing  $\Omega_{f(x_0)} f$ .

Then

$$\sum_{k \geq 0} (\cos \theta_k)^2 \cdot \|\nabla f(x_k)\|^2 < \infty,$$

where  $\theta_k := \sphericalangle(-\nabla f(x_k), p_k)$

### Result

$$\cos \theta_k \geq \delta > 0 \quad \Rightarrow \quad \lim_{k \rightarrow \infty} \|\nabla f(x_k)\| \rightarrow 0$$

## Gradient Descent

Given a (continuously) differentiable objective  $f: \mathbb{R}^n \rightarrow \mathbb{R}$ .

We have information concerning the slope(s) of  $f$ .

### Algorithm: Gradient Descent

Use a descent method (with or without line search).

Use the search direction

$$\forall k : p_k := -\nabla f(x_k)$$

Locally (i.e.  $\bar{\alpha} \rightarrow 0$ ), this is the best possible choice



# Gradient Descent: Convergence

Gradient Descent for quadratic test problem

$$f(x) := \frac{1}{2} x^\top \cdot Q \cdot x - b^\top \cdot x$$

with positive definite matrix  $Q$ .

## Gradient Descent with Exact Line Search

Update formula

$$x_{k+1} = x_k - \frac{\nabla f(x_k)^\top \cdot \nabla f(x_k)}{\nabla f(x_k)^\top \cdot Q \cdot \nabla f(x_k)} \cdot \nabla f(x_k)$$

## Theorem

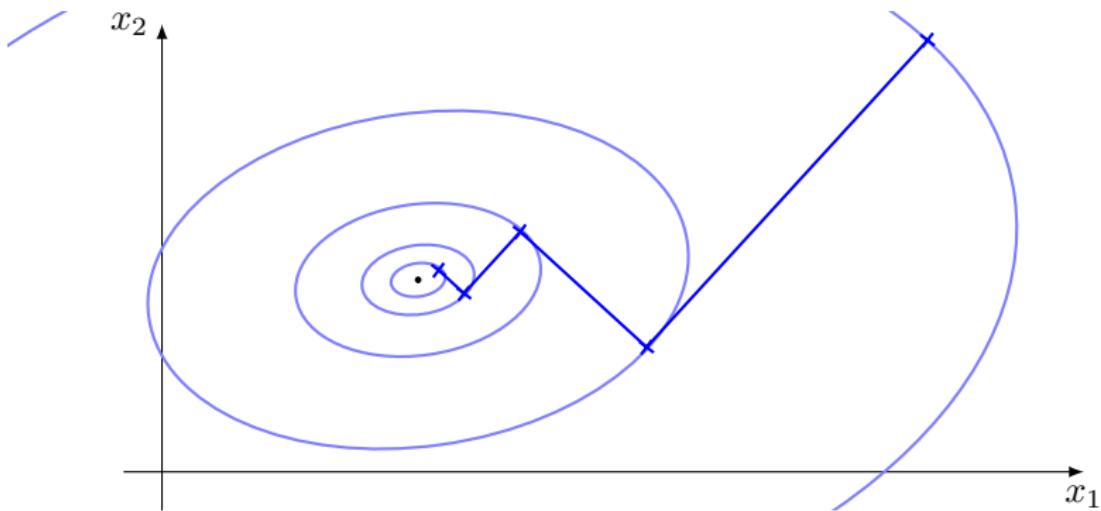
When gradient descent with exact line search is applied to  $f$  as above, then

$$\|x_{k+1} - x^*\|_Q^2 \leq \left( \frac{\lambda_n - \lambda_1}{\lambda_n + \lambda_1} \right)^2 \cdot \|x_k - x^*\|_Q^2 ,$$

where  $0 < \lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_n$  are the eigenvalues of  $Q$ . (linear convergence)

Similar result for nonlinear case. ( $Q \leftrightarrow \nabla^2 f(x^*)$ ,  $(\cdot)^2 \leftrightarrow r^2$  with  $\kappa(Q) < r < 1$ )

## Gradient Descent: Convergence



Perfect (convergence in one step) if all level curves equally spaced along all axes.

# Newton's Method



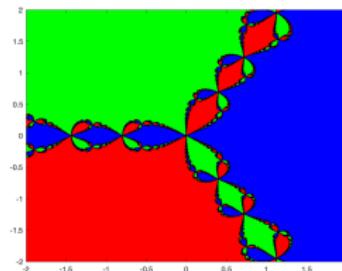
Sir Isaac Newton, 1643 - 1727

## Motivation

- Instant method for  $f(x) = x^\top \cdot Q \cdot x - b^\top \cdot x$  with s.p.d. matrix  $Q$
- Solve  $\nabla f = \mathbf{0}$  with the 'classical' Newton's method
- (Sequential) approximate of  $f$  by quadratic functions

## Global Convergence Properties

- Oscillations
- Divergence
- Stuck in local minima
- Often not even a descent direction
- Not defined if  $\nabla^2 f(x_k)$  singular



Source: By Godfrey Kneller - <http://www.phys.uu.nl/vgent/astrology/images/newton1689.jpg>], Public Domain,

## Newton's Method

### Newton's Method

Apply the descent method algorithm with the update

$$x_{k+1} := x_k - \underbrace{\alpha_k}_{\stackrel{=1}{\text{or } \rightarrow 1}} \cdot \underbrace{\left(\nabla^2 f(x_k)\right)^{-1}}_{\substack{\text{not always} \\ \text{s.p.d.}}} \cdot \nabla f(x_k)$$

### (Additional) Costs:

- approximation/calculation of Hessian  $\nabla^2 f(x_k)$ , storage thereof
- Inversion of  $\nabla^2 f(x_k)$  (never invert(!))

### Why all this trouble?

### Theorem

If  $f$  is twice differentiable with a L.-continuous Hessian and  $\nabla^2 f(x^*)$  is s.p.d. in a local minimum  $x^*$  and  $x_0$  is **sufficiently close** to  $x^*$ , then Newton's method **converges quadratically** in  $x$  and  $\nabla f$  towards  $x^*$ .

## Modifications of Newton's Method \*different names in literature

### Modified Newton's Method

Replace

$$\nabla^2 f(x_k) \rightsquigarrow \nabla^2 f(x_{0/m}) \quad (m < k)$$

loose quadratic convergence, fewer evaluations/inversions

### Simplified Newton

Replace Hessian by (good and easy to invert) approximation

$$\nabla^2 f(x_k) \rightsquigarrow Q_k \approx \nabla^2 f(x_k)$$

### Damped Newton

Use a reduced step length  $\alpha_k < 1$  in the Newton updates.

loose quadratic convergence, (often:) increase domain of convergence

### In Practice

Newton's method only when Hessian cheap and/or to finalize iteration.

## (Linear) Conjugated Gradient Method (I)

**Back to quadratic problem:**

$$f(x) = \frac{1}{2} \cdot x^\top \cdot A \cdot x - b^\top \cdot x$$

with a s.p.d. matrix  $A \in \mathbb{R}^{n \times n}$ .

**Alternatively:** Solve the (large) linear system  $A \cdot x = b$ . with the residual  $r(x) := A \cdot x - b = \nabla f(x)$

### Conjugate Direction Methods

Use a set of conjugate directions  $\{p_0, p_1, \dots, p_{n-1}\}$  (always lin. indep.) as search directions in the descent algorithm. i.e.  $p_i^\top \cdot A \cdot p_j = 0$ , ( $i \neq j$ )

- Exact line search like for the gradient method possible
- Convergence after  $n$  steps guaranteed
- After  $k$  steps: Minimization of  $f$  over a  $k$ -dimensional subspace (a so-called Krylov space ( $\rightarrow$  Cayley-Hamilton)))

## (Linear) Conjugated Gradient Method (II)

CG method: Basic idea

- Use only previous vector  $p_{k-1}$  and residual  $r(x_k)$  to construct next conjugate direction  $p_k$

$$p_k = -r(x_k) + \beta_k \cdot p_{k-1} \quad (\beta_k \text{ uniquely det. from being conj. direction})$$

Linear CG method

$$r_0 = Ax_0 - b, p_0 = -r_0$$

**While**  $r_k \neq 0$ :

$$\alpha_k = -\frac{r_k^\top \cdot p_k}{p_k^\top \cdot A \cdot p_k}$$

$$x_{k+1} = x_k + \alpha_k \cdot p_k$$

$$r_{k+1} = A \cdot x_{k+1} - b$$

$$\beta_{k+1} = \frac{r_{k+1}^\top \cdot A \cdot p_k}{p_k^\top \cdot A \cdot p_k}$$

$$p_{k+1} = -r_{k+1} + \beta_{k+1} \cdot p_k$$

## Convergence and Preconditioned CG

### Theorem

If  $A$  has only  $r$  distinct eigenvalues, then the (linear) conjugated gradient method terminates at the solution after  $r$  steps.

### Theorem

The (theoretical) convergence rate is similar to the gradient method

$$\|x_{k+1} - x^*\|_A^2 \leq \left( \frac{\lambda_{n-k} - \lambda_1}{\lambda_{n-k} + \lambda_1} \right)^2 \cdot \|x_0 - x^*\|_A^2$$

### PCG

Improve convergence behavior by linearly transforming the search space

$$\hat{x} := C \cdot x$$

$$\hat{f}(\hat{x}) = \frac{1}{2} \cdot \hat{x}^\top \cdot (C^{-\top} A C^{-1}) \hat{x} - (C^{-\top} b)^\top \cdot \hat{x}$$

**NB:** Extension (and convergence analysis) to the nonlinear case (relatively) easy.

## Least Squares Problems

Form of the objective function:

$$f(x) = \frac{1}{2} \sum_{j=1}^m r_j^2(x)$$

where each  $r_j: \mathbb{R}^n \rightarrow \mathbb{R}$  is a real valued smooth function called *residual*. We assume that  $m \geq n$ .

### Remark

- Least-squares problems arise in many areas of applications, and may in fact be the largest source of unconstrained optimization problems
- measure the discrepancy between the model and the observed behavior of the system.
- by minimizing  $f$ , we select values for the parameters that best match the model to the data.
- we show how to devise efficient, robust minimization algorithms by exploiting the special structure of the function  $f$  and its derivatives.

## Least Squares Problems (cont.)

Residual vector  $r: \mathbb{R}^n \rightarrow \mathbb{R}^m$

$$r(x) := (r_1(x), \dots, r_m(x))^\top$$

Then  $f = \frac{1}{2} \|r(x)\|_2^2$ . The *Jacobian* ( $m \times n$ -matrix) is defined as

$$J(x) := \left[ \frac{\partial r_j}{\partial x_i} \right]_{j=1, \dots, m, i=1, \dots, n} = \begin{pmatrix} \nabla r_1(x)^\top \\ \vdots \\ \nabla r_m(x)^\top \end{pmatrix}.$$

$$\nabla f(x) = \sum_{j=1}^m r_j(x) \nabla r_j(x) = J(x)^\top r(x),$$

$$\begin{aligned} \nabla^2 f(x) &= \underbrace{\sum_{j=1}^m \nabla r_j(x) \nabla r_j(x)^\top}_{\text{"for free"}} + \sum_{j=1}^m r_j(x) \nabla^2 r_j(x)^\top \\ &= J(x)^\top J(x) + \sum_{j=1}^m r_j(x) \nabla^2 r_j(x). \end{aligned} \tag{1}$$

# Gauss Newton Methods

## The Gauss-Newton Method

The standard Newton equations  $\nabla^2 f(x_k)p = -\nabla f(x_k)$  are replaced by the following system to obtain the search direction  $p_k^{\text{GN}}$ :

$$J_k^\top J_k p_k^{\text{GN}} = -J_k^\top r_k.$$

### Advantages

Using the approximation  $\nabla^2 f(x_k) \approx J_k^\top J_k$  saves the trouble of computing the individual residual Hessians  $\nabla^2 r_j$  from (1). Often, the first term in (1) dominates the second one, such that  $J_k^\top J_k$  is a good approximation of  $\nabla^2 f(x_k)$ . In practice, many least-squares problems have small residuals at the solution, leading to rapid local convergence of Gauss-Newton.

### Convergence

In theory: (Linear) convergence under mild regularity and full-rank condition on system matrices.

In practice: Good global and fast local convergence

# What Next?

The screenshot shows a web browser displaying the neos-guide.org website. The URL in the address bar is <https://neos-guide.org/content/algorithms-by-type>. The page title is "neos Guide". Below the title, there is a navigation menu with links to "Case Studies", "Optimization Guide", "Server Resources", "Server Information", and "NEOS Server". The main content area has a blue header "Optimization Guide" and a section title "Algorithms by Problem Type". A text block below the title states: "Included below is a list of algorithms, organized by optimization problem type, that have a linked page. To learn more about these algorithms, see [Algorithms](#)". There are two sections of bulleted lists: "Unconstrained Optimization" and "Related Problems".

**Optimization Guide**

## Algorithms by Problem Type

Included below is a list of algorithms, organized by optimization problem type, that have a linked page. To learn more about these algorithms, see [Algorithms](#).

### Unconstrained Optimization

- Line Search Methods
- Trust-Region Methods
- Truncated Newton Methods
- Difference Approximations
- Quasi-Newton Methods
- Nonlinear Conjugate Gradient Method
- Nonlinear Simplex Method

### Related Problems

- Nonlinear Least-Squares Problems
  - The Gauss Newton Method
  - The Levenberg-Marquardt Method
  - Hybrid Methods
  - Large-Scale Methods
  - Nonlinear Least-Squares Problems with Constraints
- Nonlinear Equations
  - Trust-Region and Line Search Methods
  - Broyden's Method
  - Tensor Methods
  - Homotopy Methods

### Constrained Optimization